

# Machine Learning Odyssey Part 6

Feb 18, 2024

Seungeun Lee

# So far...

- Tabular data (data preprocessing, data imbalance, ...)
- AutoML
- XGBoost (CNN of machine learning), LightGBM
- TabNet
- Random Forest
- ...

**① | Tabular Data**

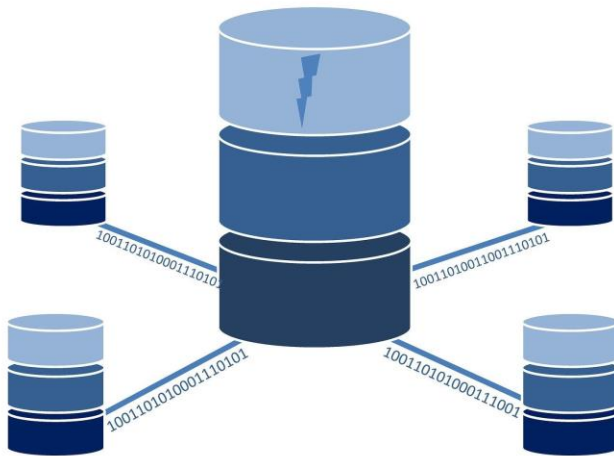
**② | SVM**

**③ | AutoML**

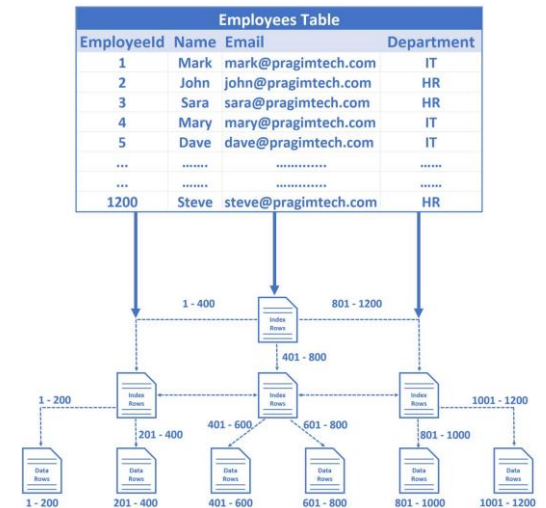
**④ | Hyperparameter Tuning**

# Tabular data

- Tabular data  $\in$  Structured Data
- Data extracted from a database (DB), represented in a tabular format with rows and columns



DataBase



Data

x		target	o		x		o		o			
PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S
...	...	...	...	...	...	...	...	...	...	...	...	...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.0000	NaN	S
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.0000	B42	S
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607	23.4500	NaN	S
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.0000	C148	C
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.7500	NaN	Q
891 rows × 12 columns												

Titanic dataset (<https://inovation97.tistory.com/65>)

# Tabular vs. Image data

Tabular Data

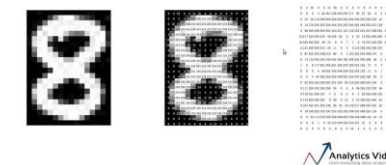
columns = attributes for those observations

Player	Minutes	Points	Rebounds	Assists
A	41	20	6	5
B	30	29	7	6
C	22	7	7	2
D	26	3	3	9
E	20	19	8	0
F	9	6	14	14
G	14	22	8	3
I	22	36	0	9
J	34	8	1	3

Rows = observations

<https://www.kaggle.com/competitions/kdd2021/tabular1.jpg>

VS.



- Image data  $\in$  Unstructured Data
- While there is a lack of overall structural differences btw tabular and image data, what are the distinctions?

## 1. Difference in Storage Methods – Database vs. Cloud

## 2. Level of Feature Extraction – Processed vs. Raw

- The second point reflects the **difference in surrounding values**:

In tabular data, adjacent values are less correlated, while in image data, adjacent pixels are highly correlated (leading to methods like CNN and the concept of the receptive field).

```
select * from users
```

users 1 X

 `select * from users` |  Enter a SQL expression to filter results (use Ctrl+Space)

	 id	 first_name	 last_name	 email
1	5,316	Steven	Woods	stevenwoods@e
2	13,446	Rebecca	Santos	rebeccasantos@
3	1,978	William	Ortiz	williamortiz@ex
4	4,834	Jennifer	Patterson	jenniferpatterson
5	37,153	Victor	Martin	victormartin@ex
6	49,725	Christina	Wood	christinawood@

# Multimodal Era

## Cloud

- Google Cloud Platform (GCP)
- Amazon Web Services (AWS)
- Samsung Cloud Platform (SCP)



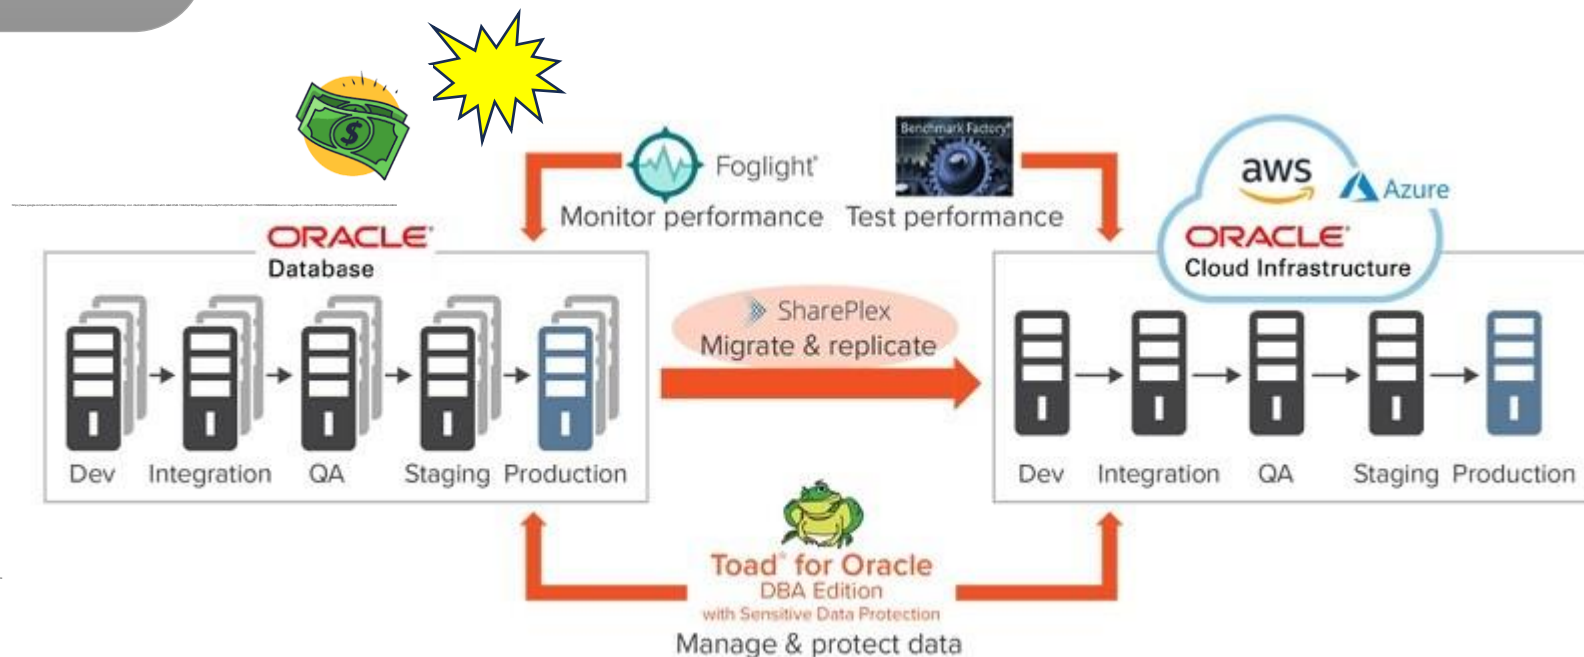
Meow



고양이 (Korean), gato (Español)

Cat

Tabular data related to feed cats





**① | Tabular Data**

**② | SVM**

**③ | AutoML**

**④ | Hyperparameter Tuning**

# Multimodal Era



Meow



고양이 (Korean), gato (Español)

Cat

Tabular data related to feed cats



Image data: CNN, Vision Transformer



Audio data: Wav2vec 2.0, AASIST



Text data (Korean): KoBART, KoBERT / Español



Text data (English): BERT, GPT, ELECTRA



?????

# Machine Learning for Tabular data

- XGBoost (CNN of machine learning), LightGBM
  - Random Forest
  - **Deep Learning approaches**
    - TabNet, STUNT, Trompt, ...
    - Does not work well, though ... (overparametrized, nbhd values, ...)
- > ?????

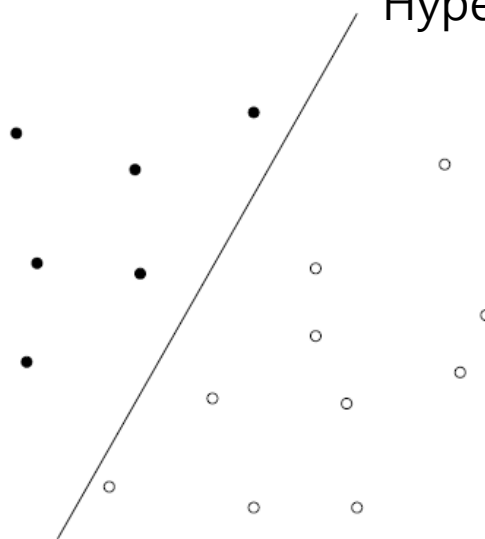
# Support Vector Machine [SVM]

- Finding the optimal hyperplane
- Binary Classification

separate two sets of points  $\{x_1, \dots, x_N\}$ ,  $\{y_1, \dots, y_M\}$  by a hyperplane:

$$a^T x_i + b > 0, \quad i = 1, \dots, N, \quad a^T y_i + b < 0, \quad i = 1, \dots, M$$

Hyperplane, linear classifier



homogeneous in  $a, b$ , hence equivalent to For an arbitrarily small  $\varepsilon > 0$ ,

$$a^T x_i + b \geq \varepsilon, \quad i = 1, \dots, N, \quad a^T y_i + b \leq -\varepsilon, \quad i = 1, \dots, M$$

a set of linear inequalities in  $a, b$

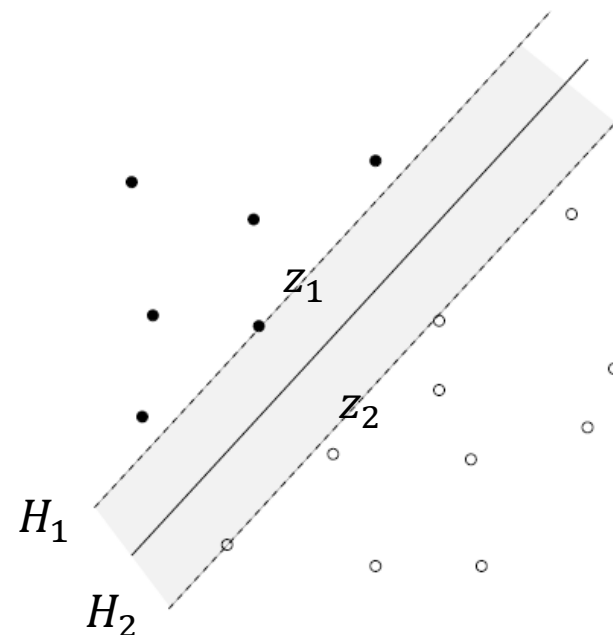
# Robust linear classifier

(Euclidean) distance between hyperplanes

$$\mathcal{H}_1 = \{z \mid a^T z + b = 1\}$$

$$\mathcal{H}_2 = \{z \mid a^T z + b = -1\}$$

Maximize the distance btw  $z_1$  and  $z_2$   
is  $\text{dist}(\mathcal{H}_1, \mathcal{H}_2) = 2/\|a\|_2$



to separate two sets of points by maximum margin,

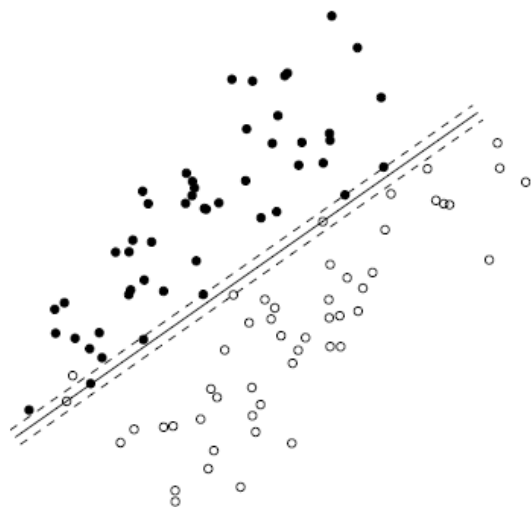
$$\begin{aligned} &\text{minimize} && (1/2)\|a\|_2 && \text{Minimize its reciprocal} \\ &\text{subject to} && a^T x_i + b \geq 1, && i = 1, \dots, N \\ &&& a^T y_i + b \leq -1, && i = 1, \dots, M \end{aligned} \tag{1}$$

(after squaring objective) a QP in  $a, b$

# Approximate linear classifier

$$\begin{array}{ll}\text{minimize} & \mathbf{1}^T u + \mathbf{1}^T v \\ \text{subject to} & a^T x_i + b \geq 1 - u_i, \quad i = 1, \dots, N \\ & a^T y_i + b \leq -1 + v_i, \quad i = 1, \dots, M \\ & u \succeq 0, \quad v \succeq 0\end{array}$$

- an LP in  $a, b, u, v$
- at optimum,  $u_i = \max\{0, 1 - a^T x_i - b\}$ ,  $v_i = \max\{0, 1 + a^T y_i + b\}$
- can be interpreted as a heuristic for minimizing #misclassified points



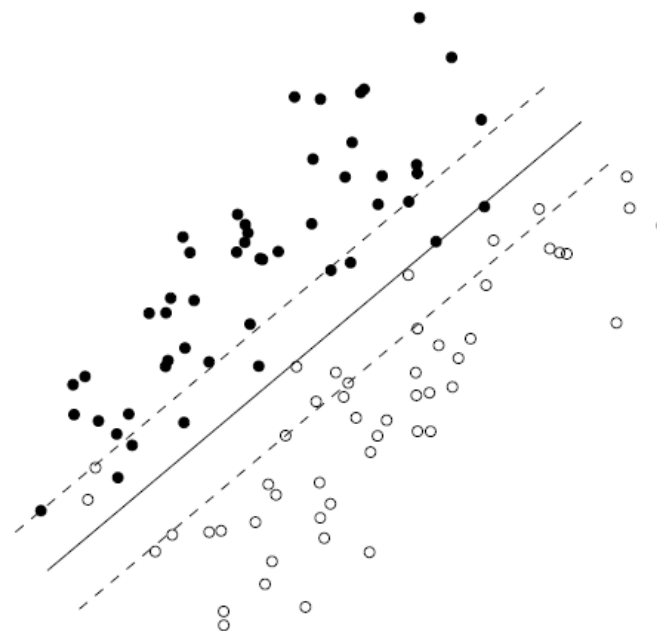
- minimizing the number of misclassification is hard, but this approach is good approximation

# Support vector machine (SVM)

$$\begin{aligned} & \text{minimize} && \|a\|_2 + \gamma(\mathbf{1}^T u + \mathbf{1}^T v) \\ & \text{subject to} && a^T x_i + b \geq 1 - u_i, \quad i = 1, \dots, N \\ & && a^T y_i + b \leq -1 + v_i, \quad i = 1, \dots, M \\ & && u \succeq 0, \quad v \succeq 0 \end{aligned} \quad \gamma: \text{hyperparameter}$$

produces point on trade-off curve between inverse of margin  $2/\|a\|_2$  and classification error, measured by total slack  $\mathbf{1}^T u + \mathbf{1}^T v$

same example as previous page,  
with  $\gamma = 0.1$ :





# Support Vector Machine [SVM]

- Multi-class classification
- OvO [One vs. One]
- OvR [One vs. Rest]

e.g.) cat vs. dog vs. fig

OvO: cat – dog | cat – fig | dog – fig

OvR: cat – {dog, fig} | dog – {cat, fig} | fig – {cat, dog}

**① | Tabular Data**

**② | SVM**

**③ | AutoML**

**④ | Hyperparameter Tuning**

# AutoML

## AutoML-pycaret

- An open-source library that automates the machine learning workflow with simple code, encompassing coding, preprocessing, model selection, and parameter tuning all in one!
- **But... How to Improve Tuning?**

**① | Tabular Data**

**② | SVM**

**③ | AutoML**

**④ | Hyperparameter Tuning**

# Hyperparameter Tuning

*A modern nightmare of data scientists...*

(1) Grid Search – computes all possible cases

<https://dacon.io/codeshare/4646>

```
from sklearn.model_selection import GridSearchCV

grid_search = {'criterion': ['entropy', 'gini'],
               'max_depth': [2],
               'max_features': ['auto', 'sqrt'],
               'min_samples_leaf': [4, 6, 8],
               'min_samples_split': [5, 7, 10],
               'n_estimators': [20]}

clf = RandomForestClassifier()
model = GridSearchCV(estimator = clf, param_grid = grid_search, cv = 4, verbose= 5, n_jobs = -1)
model.fit(X_Train, Y_Train)
```

# Hyperparameter Tuning

## (2) Random Search – randomly selects cases

<https://dacon.io/codeshare/4646>

```
from sklearn.model_selection import RandomizedSearchCV

random_search = {'criterion': ['entropy', 'gini'],
                  'max_depth': [2],
                  'max_features': ['auto', 'sqrt'],
                  'min_samples_leaf': [4, 6, 8],
                  'min_samples_split': [5, 7, 10],
                  'n_estimators': [20]}

clf = RandomForestClassifier()
model = RandomizedSearchCV(estimator = clf, param_distributions = random_search, n_iter = 10,
                           cv = 4, verbose= 1, random_state= 101, n_jobs = -1)
model.fit(X_Train, Y_Train)
```

# Hyperparameter Tuning

## (3) HyperOpt

<https://dacon.io/codeshare/4646>

```
from hyperopt import hp, fmin, tpe, STATUS_OK, Trials

space = {'criterion': hp.choice('criterion', ['entropy', 'gini']),
        'max_depth': hp.quniform('max_depth', 10, 12, 10),
        'max_features': hp.choice('max_features', ['auto', 'sqrt', 'log2',
None]),
        'min_samples_leaf': hp.uniform('min_samples_leaf', 0, 0.5),
        'min_samples_split': hp.uniform('min_samples_split', 0, 1),
        'n_estimators': hp.choice('n_estimators', [10, 50])
}

def objective(space):
    hopt = RandomForestClassifier(criterion = space['criterion'],
                                max_depth = space['max_depth'],
                                max_features = space['max_features'],
                                min_samples_leaf = space['min_samples_leaf'],
                                min_samples_split = space['min_samples_split'],
                                n_estimators = space['n_estimators'],

    )
```

```
accuracy = cross_val_score(hopt, X_Train, Y_Train, cv = 4).mean()
return {'loss': -accuracy, 'status': STATUS_OK }

trials = Trials()
best = fmin(fn= objective,
           space= space,
           algo= tpe.suggest,
           max_evals = 20,
           trials= trials
)

# optimal solution
best
```

# HyperOpt

<https://towardsdatascience.com/hyperopt-demystified-3e14006eb6fa>

- Tree-based Parzen Estimators (TPE) – optimizes a user-defined objective function

1. Train a model with several sets of randomly-selected hyperparameters, returning objective function values.
2. Split our observed objective function values into “good” and “bad” groups, according to some threshold gamma ( $\gamma$ ).
3. Calculate the “promisingness” score, which is just  $P(x/good) / P(x/bad)$ .
4. Determine the hyperparameters that maximize promisingness score via mixture models.
5. Fit our model using the hyperparameters from step 4.
6. Repeat steps 2–5 until a stopping criteria.



# HyperOpt

<https://towardsdatascience.com/hyperopt-demystified-3e14006eb6fa>

- Tree-based Parzen Estimators (TPE) ∈ Sequential Model-Based Optimization (SMBO) algorithm

```
from hyperopt import hp, fmin, tpe, STATUS_OK, Trials

space = {'criterion': hp.choice('criterion', ['entropy', 'gini']),
        'max_depth': hp.quniform('max_depth', 10, 12, 10),
        'max_features': hp.choice('max_features', ['auto', 'sqrt', 'log2',
        None]),
        'min_samples_leaf': hp.uniform('min_samples_leaf', 0, 0.5),
        'min_samples_split': hp.uniform('min_samples_split', 0, 1),
        'n_estimators': hp.choice('n_estimators', [10, 50])
}

def objective(space):
    hopt = RandomForestClassifier(criterion = space['criterion'],
                                max_depth = space['max_depth'],
                                max_features = space['max_features'],
                                min_samples_leaf = space['min_samples_leaf'],
                                min_samples_split = space['min_samples_split'],
                                n_estimators = space['n_estimators'],

    )
```

constraints

objective function (Accuracy, AUC, RMSE, ...)

# HyperOpt

<https://towardsdatascience.com/hyperopt-demystified-3e14006eb6fa>

- Bayesian optimization: a sequential algorithm that finds points in hyperspace with a high probability of being “successful” according to an objective function
- Clever tricks? modeling  $P(x|y)$  instead of  $P(y|x)$ :  
$$p(x|y) = \frac{p(y|x)*p(x)}{p(y)}$$
- probability of an objective function value (y), given hyperparameters (x) -> b.o.
- **$P(x|y)$  :**

=> TPE tries to find the best objective function values, then determine the associated hyperparameters

# HyperOpt

<https://towardsdatascience.com/hyperopt-demystified-3e14006eb6fa>

- TPE splits our observed data points into two groups

(1)  $g(x)$ : good

(2)  $\ell(x)$ : bad

$$p(x|y) = \begin{cases} \ell(x) & \text{if } y < y^* \text{ Bad Objective Function Vals} \\ g(x) & \text{if } y \geq y^* \text{ Good Objective Function Vals} \end{cases}$$

# HyperOpt

<https://towardsdatascience.com/hyperopt-demystified-3e14006eb6fa>

- “Promisingness” score

$$P = \frac{g(x)}{l(x)} = \frac{P(x|good)}{P(x|bad)}$$

- Numerator: the probability of observing a set of hyperparameters (x), given the corresponding objective function value is “good.”
- Denominator: the probability of observing a set of hyperparameters (x), given the corresponding objective function value is “bad.”

=> The bigger the “promisingness” value, the more likely that the hyperparameters x will produce a “good” objective function

# HyperOpt

<https://towardsdatascience.com/hyperopt-demystified-3e14006eb6fa>

- “Promisingness” score

$$P = \frac{g(x)}{l(x)} = \frac{P(x|good)}{P(x|bad)}$$

- Numerator: the probability of observing a set of hyperparameters (x), given the corresponding objective function value is “good.”
- Denominator: the probability of observing a set of hyperparameters (x), given the corresponding objective function value is “bad.”

=> The bigger the “promisingness” value, the more likely that the hyperparameters x will produce a “good” objective function

# HyperOpt

<https://towardsdatascience.com/hyperopt-demystified-3e14006eb6fa>

- Bayesian hyperparameter optimization
- develop a probabilistic distribution of the hyperparameter search space, using an acquisition function, such as expected improvement, to transform the hyperspace to be more “searchable.”
- Then leverages optimization algorithm, such as stochastic gradient descent, to find a the hyperparameters that maximize our acquisition function.

A function that suggests what values to try next based on the current probabilistic estimates of the objective function.

=> “Promisingness” score acts as an acquisition function and is proportional to the Expected Improvement (EI)

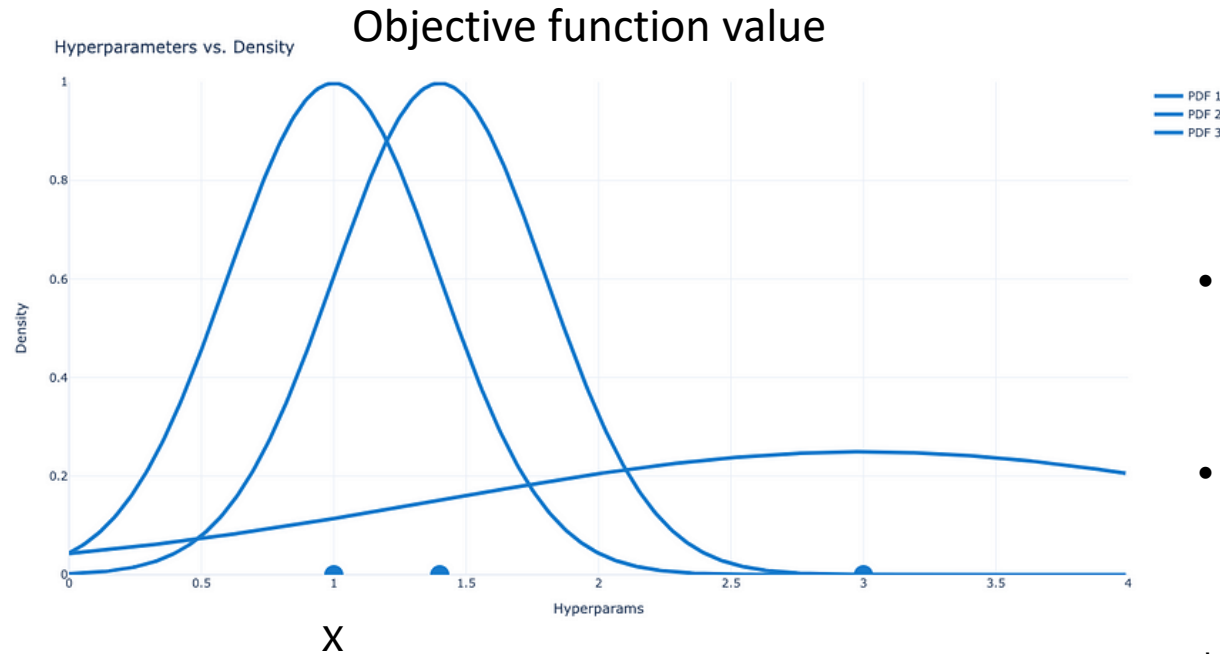
# HyperOpt

<https://towardsdatascience.com/hyperopt-demystified-3e14006eb6fa>

- “Promisingness” score is measured by the mixture model
- take multiple probability distributions and put them together using a linear combination
- Generally
  - \* TPE also support categorical variables which traditional Bayesian optimization does not.
- (1) Distribution types: categorical -> re-weighted categorical distribution / numeric-> Gaussian (i.e. normal) or uniform distribution.
- (2) Iterate over each point and insert a distribution at that point.
- (3) Sum the mass of all distributions to get a probability density estimate.
- Note that this process is run individually for both sets  $l(x)$  and  $g(x)$ .

# HyperOpt

<https://towardsdatascience.com/hyperopt-demystified-3e14006eb6fa>



- If points are close together -> the standard deviation will be small -> the distribution will be very tall
- if points are spread apart, the distribution will be flat

\* Truncated Gaussian

- $g(x) = y$
- $x$ : hyperparameters,  $y$ : objective function value
- $\sigma$ : the standard deviation
- The distance to the closest neighboring point



# HyperOpt

<https://towardsdatascience.com/hyperopt-demystified-3e14006eb6fa>

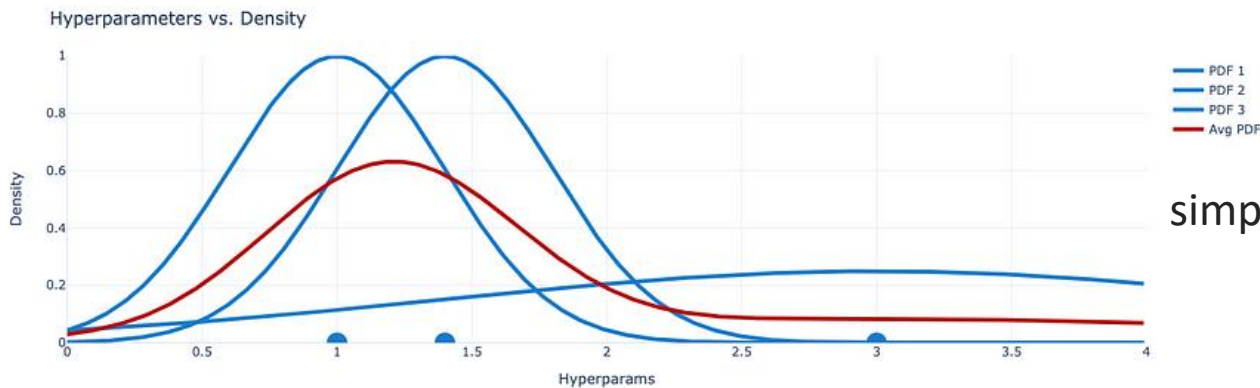
- Finally, let's determine the next point to explore! – acquisition function

(1) acquired objective function observations

(2) defined our “promisingness” formula

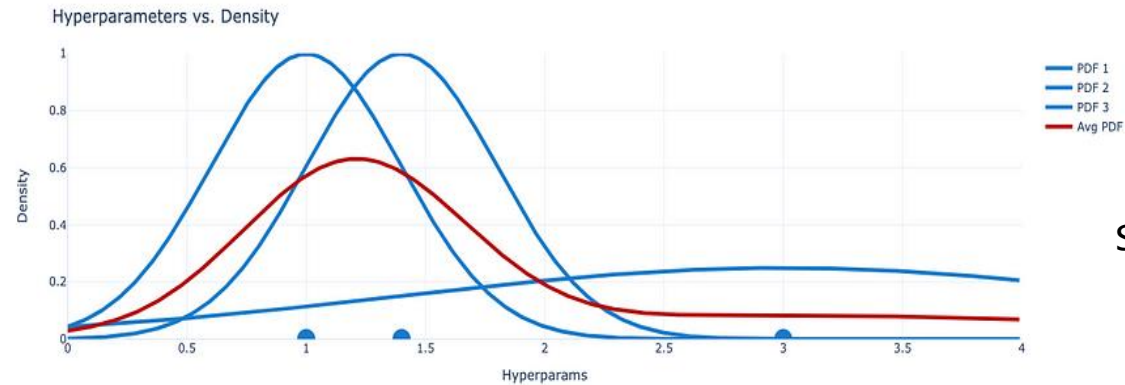
(3) created a probability density estimate via mixture models based on prior values

- average probability density function (PDF) for both  $g(x)$  and  $l(x)$



simply the sum of all PDFs divided by the number of PDFs

# HyperOpt



- Using the average PDF, we can get the probability of any hyperparameter value ( $x$ ) being in  $g(x)$  or  $l(x)$ .
- hyperparameter value ( $x$ ) of 3.9 or 0.05 -> unlikely to belong to the “good” set
- 1 to 1.2 -> seems to be very likely to belong to the “good” set
- Do the same thing to  $l(x)$
- Objective: maximize  $g(x) / l(x)$
- **promising points should be located where  $g(x)$  is high, and  $l(x)$  is low**
- With these probability distributions, sample cases from tree-structured hyperparameters (TPE..!!) and find the set of hyperparameters that maximize “promisingness”

# HyperOpt

- Algorithms for Hyper-Parameter Optimization
- [https://papers.nips.cc/paper\\_files/paper/2011/hash/86e8f7ab32cfd12577bc2619bc635690-Abstract.html](https://papers.nips.cc/paper_files/paper/2011/hash/86e8f7ab32cfd12577bc2619bc635690-Abstract.html)

# Hyperparameter Tuning

Grid Search, Random Search, HyperOpt -> “methods” for hyperparameter tuning

(4) **Optuna** – computes all possible cases within a given range -> works like AutoML

<https://dacon.io/codeshare/4646>

```
def objective(trial):
    iris = sklearn.datasets.load_iris()
    x, y = iris.data, iris.target

    classifier_name = trial.suggest_categorical('classifier', ['SVC', 'RandomForest'])

    if classifier_name == 'SVC':
        svc_c = trial.suggest_loguniform('svc_c', 1e-10, 1e10)
        classifier_obj = sklearn.svm.SVC(C=svc_c, gamma='auto')
    else:
        rf_max_depth = int(trial.suggest_loguniform('rf_max_depth', 2, 32))
        classifier_obj = sklearn.ensemble.RandomForestClassifier(max_depth=rf_max_depth, n_estimators=10)

    accuracy = cross_val_score(classifier_obj, x, y, cv = 4).mean()
    return accuracy

study = optuna.create_study(direction='maximize')
study.optimize(objective, n_trials=100)
print(study.best_trial.params)
```

# Optuna

[https://jhtobigs.oopy.io/optuna\\_tutorial](https://jhtobigs.oopy.io/optuna_tutorial)

## SOTA Method

- **Study-Trial method**
- **Study:** A session for optimizing the objective function, composed of multiple trials.
- **Default:** TPE (Tree-structured Parzen Estimator) from HyperOpt; other methods include Random Search and Grid Search.
- The best method is selected based on results.
- **Pruning:** Early termination of trials that yield poor results.